

CS 33

Machine Programming (2)

Most of the slides in this lecture are either from or adapted from slides provided by the authors of the textbook “Computer Systems: A Programmer’s Perspective,” 2nd Edition and are provided from the website of Carnegie-Mellon University, course 15-213, taught by Randy Bryant and David O’Hallaron in Fall 2010. These slides are indicated “Supplied by CMU” in the notes section of the slides.

Processor State (x86-64, Partial)

%rax	%eax	%r8	%r8d
%rbx	%ebx	%r9	%r9d
%rcx	%ecx	%r10	%r10d
%rdx	%edx	%r11	%r11d
%rsi	%esi	%r12	%r12d
%rdi	%edi	%r13	%r13d
%rsp	%esp	%r14	%r14d
%rbp	%ebp	%r15	%r15d
%rip		CF ZF SF OF	
		condition codes	

Condition Codes (Implicit Setting)

- **Single-bit registers**

CF carry flag (for unsigned)

SF sign flag (for signed)

ZF zero flag

OF overflow flag (for signed)

- **Implicitly set (think of it as side effect) by arithmetic operations**

example: *addl/addq Src, Dest* \leftrightarrow *t = a+b*

CF set if carry out from most significant bit (unsigned overflow)

ZF set if *t == 0*

SF set if *t < 0* (as signed)

OF set if two's-complement (signed) overflow

(a>0 && b>0 && t<0) || (a<0 && b<0 && t>=0)

- **Not set by *leal* instruction**

Condition Codes (Explicit Setting: Compare)

- **Explicit setting by compare instruction**

`cmpl/cmpq src2, src1`

`cmpl b, a` like computing `a-b` without setting destination

CF set if carry out from most significant bit (used for unsigned comparisons)

ZF set if `a == b`

SF set if `(a-b) < 0` (as signed)

OF set if two's-complement (signed) overflow

`(a>0 && b<0 && (a-b)<0) || (a<0 && b>0 && (a-b)>0)`

Condition Codes (Explicit Setting: Test)

- **Explicit setting by test instruction**

`testl/testq src2, src1`

`testl b, a` like computing `a&b` without setting destination

- sets condition codes based on value of `Src1` & `Src2`
- useful to have one of the operands be a mask

ZF set when `a&b == 0`

SF set when `a&b < 0`

Reading Condition Codes

- **SetX instructions**
 - set single byte based on combinations of condition codes

SetX	Condition	Description
sete	ZF	Equal / Zero
setne	$\sim ZF$	Not Equal / Not Zero
sets	SF	Negative
setns	$\sim SF$	Nonnegative
setg	$\sim (SF \wedge OF) \ \& \ \sim ZF$	Greater (Signed)
setge	$\sim (SF \wedge OF)$	Greater or Equal (Signed)
setl	$(SF \wedge OF)$	Less (Signed)
setle	$(SF \wedge OF) \mid ZF$	Less or Equal (Signed)
seta	$\sim CF \ \& \ \sim ZF$	Above (unsigned)
setb	CF	Below (unsigned)

Supplied by CMU.

Reading Condition Codes (Cont.)

- **SetX instructions:**
 - set single byte based on combination of condition codes
- **Uses one of 8 addressable byte registers**
 - does not alter remaining 7 bytes
 - typically use `movzbl` to finish job

```
int gt (int x, int y)
{
    return x > y;
}
```

%rax	%eax	%ah	%al
------	------	-----	-----

Body

```
cmpl %esi, %edi    # compare x : y
setg %al           # %al = x > y
movzbl %al, %eax   # zero rest of %eax/%rax
```

Supplied by CMU, but converted to x86-64.

Jumping

- **jX instructions**
 - Jump to different part of code depending on condition codes

jX	Condition	Description
jmp	1	Unconditional
jbe	ZF	Equal / Zero
jne	\sim ZF	Not Equal / Not Zero
js	SF	Negative
jns	\sim SF	Nonnegative
jg	$\sim (SF \wedge OF) \ \& \ \sim ZF$	Greater (Signed)
jge	$\sim (SF \wedge OF)$	Greater or Equal (Signed)
jl	$(SF \wedge OF)$	Less (Signed)
jle	$(SF \wedge OF) \ \ ZF$	Less or Equal (Signed)
ja	$\sim CF \ \& \ \sim ZF$	Above (unsigned)
jb	CF	Below (unsigned)

Supplied by CMU.

Jumping

- **jX instructions:**
 - Jump to different

Quiz 1

What would be an appropriate description if the condition is $\sim CF$?

- a) above or equal (unsigned)
- b) not less (signed)
- c) incomparable

jX	Condition	Description
jmp	1	Unconditional
jbe	ZF	Equal / Zero
jne	$\sim ZF$	Not Equal / Not Zero
js	SF	Negative
jns	$\sim SF$	Nonnegative
jg	$\sim (SF \wedge OF) \ \& \ \sim ZF$	Greater (Signed)
jge	$\sim (SF \wedge OF)$	Greater or Equal (Signed)
jl	$(SF \wedge OF)$	Less (Signed)
jle	$(SF \wedge OF) \mid ZF$	Less or Equal (Signed)
ja	$\sim CF \ \& \ \sim ZF$	Above (unsigned)
jb	CF	Below (unsigned)

Supplied by CMU.

Conditional-Branch Example

```
int absdiff(int x, int y)
{
    int result;
    if (x > y) {
        result = x-y;
    } else {
        result = y-x;
    }
    return result;
}
```

```
absdiff:
    movl    %esi, %eax
    cmpl    %esi, %edi
    jle     .L6
    subl    %eax, %edx
    movl    %edx, %eax
    jmp     .L7
.L6:
    subl    %edx, %eax
.L7:
    ret
```

Body1

Body2a

Body2b

Supplied by CMU, but converted to x86-64.

Conditional-Branch Example (Cont.)

```
int goto_ad(int x, int y)
{
    int result;
    if (x <= y) goto Else;
    result = x-y;
    goto Exit;
Else:
    result = y-x;
Exit:
    return result;
}
```

```
absdiff:
    movl    %esi, %eax
    cmpl    %esi, %edi
    jle     .L6
    subl    %eax, %edx
    movl    %edx, %eax
    jmp     .L7
.L6:
    subl    %edx, %eax
.L7:
    ret
```

Body1

Body2a

Body2b

- C allows “goto” as means of transferring control
 - closer to machine-level programming style
- Generally considered bad coding style

Supplied by CMU, but converted to x86-64.

General Conditional-Expression Translation

C Code

```
val = Test ? Then_Expr : Else_Expr;
```

```
val = x > y ? x - y : y - x;
```

Goto Version

```
nt = !Test;  
if (nt) goto Else;  
val = Then_Expr;  
goto Done;  
Else:  
    val = Else_Expr;  
Done:  
    . . .
```

- Test is expression returning integer
 - == 0 interpreted as false
 - ≠ 0 interpreted as true
- Create separate code regions for then & else expressions
- Execute appropriate one

Using Conditional Moves

- **Conditional move instructions**

- instruction supports:

- if (Test) Dest \leftarrow Src

- supported in post-1995 x86 processors

- gcc does not always use them

- » wants to preserve compatibility with ancient processors

- » enabled for x86-64

- » use switch `-march=686` for IA32

- **Why use them?**

- branches are very disruptive to instruction flow through pipelines

- conditional moves do not require control transfer

C Code

```
val = Test
    ? Then_Expr
    : Else_Expr;
```

Goto Version

```
tval = Then_Expr;
result = Else_Expr;
t = Test;
if (t) result = tval;
return result;
```

Supplied by CMU.

Note that, as shown in the goto version, both the then_expr and the else_expr are computed before the test is done.

Conditional Move Example: x86-64

```
int absdiff(int x, int y) {  
    int result;  
    if (x > y) {  
        result = x-y;  
    } else {  
        result = y-x;  
    }  
    return result;  
}
```

```
absdiff:  
x in %edi      movl    %edi, %eax  
y in %esi      subl    %esi, %eax    # result = x-y  
               movl    %esi, %edx  
               subl    %edi, %edx    # tval = y-x  
               cmpl    %esi, %edi    # compare x:y  
               cmovle  %edx, %eax    # if <=, result = tval  
               ret
```

Bad Cases for Conditional Move

Expensive Computations

```
val = Test(x) ? Hard1(x) : Hard2(x);
```

- both values get computed
- only makes sense when computations are very simple

Risky Computations

```
val = p ? *p : 0;
```

- both values get computed
- may have undesirable effects

Computations with side effects

```
val = x > 0 ? x*=7 : x+=3;
```

- both values get computed
- must be side-effect free

“Do-While” Loop Example

C Code

```
int pcount_do(unsigned x)
{
    int result = 0;
    do {
        result += x & 0x1;
        x >>= 1;
    } while (x);
    return result;
}
```

Goto Version

```
int pcount_do(unsigned x)
{
    int result = 0;
loop:
    result += x & 0x1;
    x >>= 1;
    if (x)
        goto loop;
    return result;
}
```

- Count number of 1's in argument x (“popcount”)
- Use conditional branch either to continue looping or to exit loop

“Do-While” Loop Compilation

Goto Version

```
int pcount_do(unsigned x) {  
    int result = 0;  
loop:  
    result += x & 0x1;  
    x >>= 1;  
    if (x)  
        goto loop;  
    return result;  
}
```

Registers:

%edx	x
%eax	result

```
        movl    $0, %eax        # result = 0  
.L2:    # loop:  
        movl    %edx, %ecx  
        andl    $1, %ecx        # t = x & 1  
        addl    %ecx, %eax      # result += t  
        shrl    %edx            # x >>= 1  
        jne     .L2             # if !0, goto loop
```

Supplied by CMU.

Note that the condition codes are set as part of the execution of the shrl instruction.

General “Do-While” Translation

C Code

```
do  
    Body  
while (Test);
```

- **Body:** {
 Statement_1 ;
 Statement_2 ;
 ...
 Statement_n ;
}
- **Test returns integer**
 = 0 interpreted as false
 ≠ 0 interpreted as true

Goto Version

```
loop:  
    Body  
    if (Test)  
        goto loop
```

“While” Loop Example

C Code

```
int pcount_while(unsigned x) {  
    int result = 0;  
    while (x) {  
        result += x & 0x1;  
        x >>= 1;  
    }  
    return result;  
}
```

Goto Version

```
int pcount_do(unsigned x) {  
    int result = 0;  
    if (!x) goto done;  
loop:  
    result += x & 0x1;  
    x >>= 1;  
    if (x)  
        goto loop;  
done:  
    return result;  
}
```

- Is this code equivalent to the do-while version?
 - must jump out of loop if test fails

General “While” Translation

While version

```
while (Test)  
  Body
```



Do-While Version

```
if (!Test)  
  goto done;  
do  
  Body  
  while(Test) ;  
done:
```



Goto Version

```
if (!Test)  
  goto done;  
loop:  
  Body  
  if (Test)  
    goto loop;  
done:
```

“For” Loop Example

C Code

```
#define WSIZE 8*sizeof(int)
int pcount_for(unsigned x) {
    int i;
    int result = 0;
    for (i = 0; i < WSIZE; i++) {
        unsigned mask = 1 << i;
        result += (x & mask) != 0;
    }
    return result;
}
```

- Is this code equivalent to other versions?

“For” Loop Form

General Form

```
for (Init; Test; Update)  
    Body
```

```
for (i = 0; i < WSIZE; i++) {  
    unsigned mask = 1 << i;  
    result += (x & mask) != 0;  
}
```

Init

```
i = 0
```

Test

```
i < WSIZE
```

Update

```
i++
```

Body

```
{  
    unsigned mask = 1 << i;  
    result += (x & mask) != 0;  
}
```

“For” Loop → While Loop

For Version

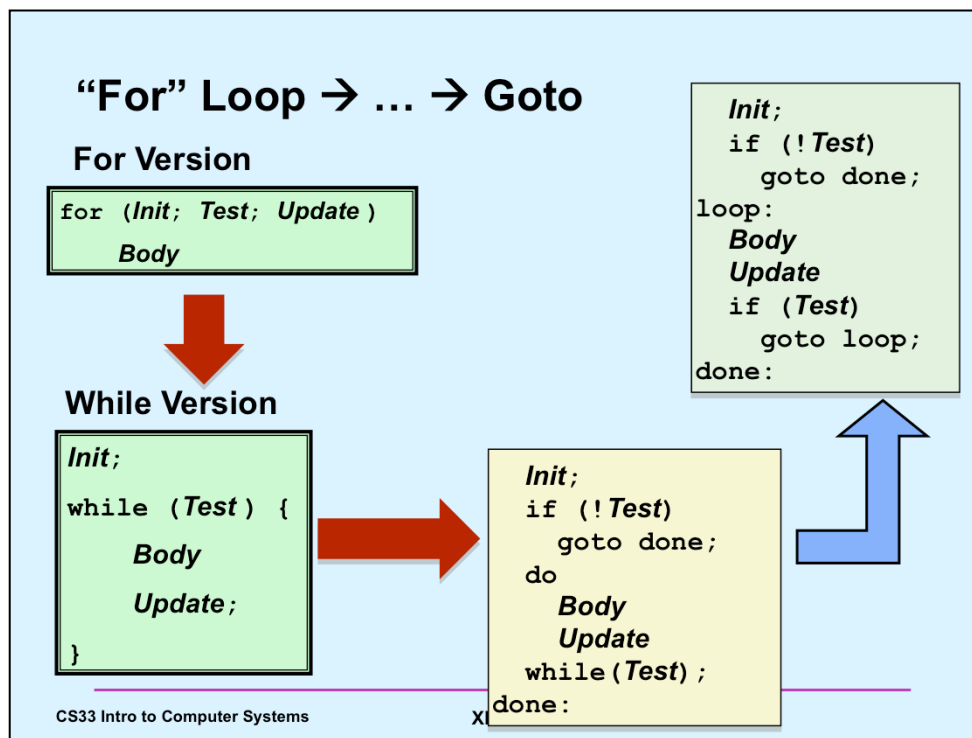
```
for (Init; Test; Update )  
    Body
```



While Version

```
Init;  
while (Test) {  
    Body  
    Update;  
}
```

Supplied by CMU.



Supplied by CMU.

“For” Loop Conversion Example

C Code

```
#define WSIZE 8*sizeof(int)
int pcount_for(unsigned x) {
    int i;
    int result = 0;
    for (i = 0; i < WSIZE; i++) {
        unsigned mask = 1 << i;
        result += (x & mask) != 0;
    }
    return result;
}
```

Initial test can be optimized away

Goto Version

```
int pcount_for_gt(unsigned x) {
    int i;
    int result = 0;   Init
    i = 0;
    if (!(i < WSIZE)) !Test
    goto done;
loop:
    {
        Body
        unsigned mask = 1 << i;
        result += (x & mask) != 0;
    }
    i++; Update
    if (i < WSIZE) Test
        goto loop;
done:
    return result;
}
```

Switch-Statement Example

```
long switch_eg
(long x, long y, long z) {
    long w = 1;
    switch(x) {
        case 1:
            w = y*z;
            break;
        case 2:
            w = y/z;
            /* Fall Through */
        case 3:
            w += z;
            break;
        case 5:
        case 6:
            w -= z;
            break;
        default:
            w = 2;
    }
    return w;
}
```

- **Multiple case labels**
 - here: 5 & 6
- **Fall-through cases**
 - here: 2
- **Missing cases**
 - here: 4

Jump-Table Structure

Switch Form

```
switch(x) {  
  case val_0:  
    Block 0  
  case val_1:  
    Block 1  
    . . .  
  case val_n-1:  
    Block n-1  
}
```

Jump Table

jtab:	Targ0
	Targ1
	Targ2
	.
	.
	.
	Targn-1

Jump Targets

Targ0:

Code Block 0

Targ1:

Code Block 1

Targ2:

Code Block 2

.
. .
.

Targn-1:

Code Block n-1

Approximate Translation

```
target = JTab[x];  
goto *target;
```

Switch-Statement Example (x86-64)

```
long switch_eg(long x, long y, long z)
{
    long w = 1;
    switch(x) {
        . . .
    }
    return w;
}
```

What range of values is covered by the default case?

Setup:

```
switch_eg:
    ...                               # Setup
    movq    %rdx, %rcx               # %rcx = z
    cmpq    $6, %rdi                 # Compare x:6
    ja      .L8                       # If unsigned > goto default
    jmp     *.L7(, %rdi, 8)           # Goto *JTab[x]
```

Note that w not initialized here

Supplied by CMU, but converted to x86-64.

Note that the *ja* in the slide causes a jump to occur if the previous comparison is interpreted as being performed on unsigned values, and the result is that *x* is greater than (above) 6. Given that *x* is declared to be a *signed* value, for what range of values of *x* will *ja* cause a jump to take place?

Note that the assembler code shown in the examples was produced by compiling the C code using *gcc* with the “-O1” flag.

Switch-Statement Example

```
long switch_eg(long x, long y, long z)
{
    long w = 1;
    switch(x) {
        . . .
    }
    return w;
}
```

Jump table

```
.section      .rodata
.align 4
.L7:
.quad        .L8 # x = 0
.quad        .L3 # x = 1
.quad        .L4 # x = 2
.quad        .L9 # x = 3
.quad        .L8 # x = 4
.quad        .L6 # x = 5
.quad        .L6 # x = 6
```

Setup:

```
switch_eg:
    ...                                # Setup
    movq    %rdx, %rcx                # %rcx = z
    cmpq    $6, %rdi                  # Compare x:6
    ja      .L8                        # If unsigned > goto default
    Indirect jump → jmp    *.L7(,%rdi,8) # Goto *JTab[x]
```

Supplied by CMU, but converted to x86-64.

Assembly-Setup Explanation

- **Table structure**

- each target requires 8 bytes
- base address at `.L7`

- **Jumping**

- direct:** `jmp .L8`
 - jump target is denoted by label `.L8`

- indirect:** `jmp *.L7(,%rdi,8)`
 - start of jump table: `.L7`
 - must scale by factor of 8 (labels have 8 bytes on x86-64)
 - fetch target from effective address `.L7 + rdi*8`
 - » only for $0 \leq x \leq 6$

Jump table

```
.section .rodata
.align 4
.L7:
.quad .L8 # x = 0
.quad .L3 # x = 1
.quad .L4 # x = 2
.quad .L9 # x = 3
.quad .L8 # x = 4
.quad .L6 # x = 5
.quad .L6 # x = 6
```

Supplied by CMU, but converted to x86-64.

Jump Table

Jump table

```
.section .rodata
.align 4
.L7:
.quad .L8 # x = 0
.quad .L3 # x = 1
.quad .L4 # x = 2
.quad .L9 # x = 3
.quad .L8 # x = 4
.quad .L6 # x = 5
.quad .L6 # x = 6
```

```
switch(x) {
case 1: // .L3
    w = y*z;
    break;
case 2: // .L4
    w = y/z;
    /* Fall Through */
case 3: // .L9
    w += z;
    break;
case 5:
case 6: // .L6
    w -= z;
    break;
default: // .L8
    w = 2;
}
```

Supplied by CMU, but converted to x86-64.

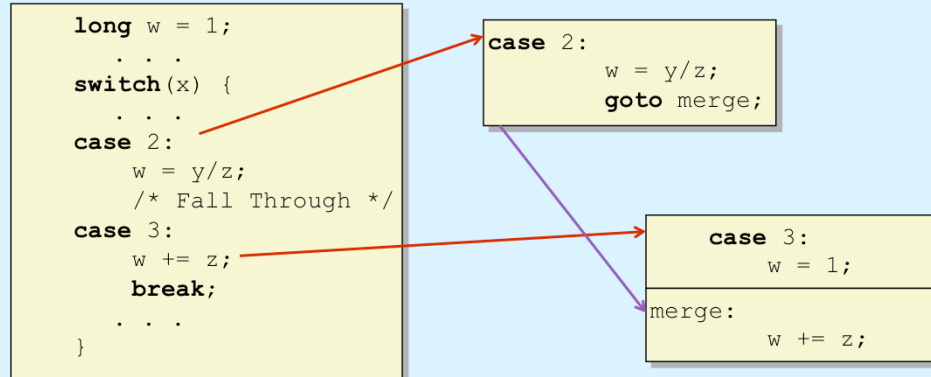
Code Blocks (Partial)

```
switch(x) {  
  case 1:      // .L3  
    w = y*z;  
    break;  
  
  . . .  
  case 5:      // .L6  
  case 6:      // .L6  
    w -= z;  
    break;  
  default:    // .L8  
    w = 2;  
}
```

```
.L3:          # x == 1  
  movl %rsi, %rax # y  
  imulq %rdx, %rax # w = y*z  
  ret  
  
.L6:          # x == 5, x == 6  
  movl $1, %eax # w = 1  
  subq %rdx, %rax # w -= z  
  ret  
  
.L8:          # Default  
  movl $2, %eax # w = 2  
  ret
```

Supplied by CMU, but converted to x86-64.

Handling Fall-Through



Supplied by CMU, but converted to x86-64.

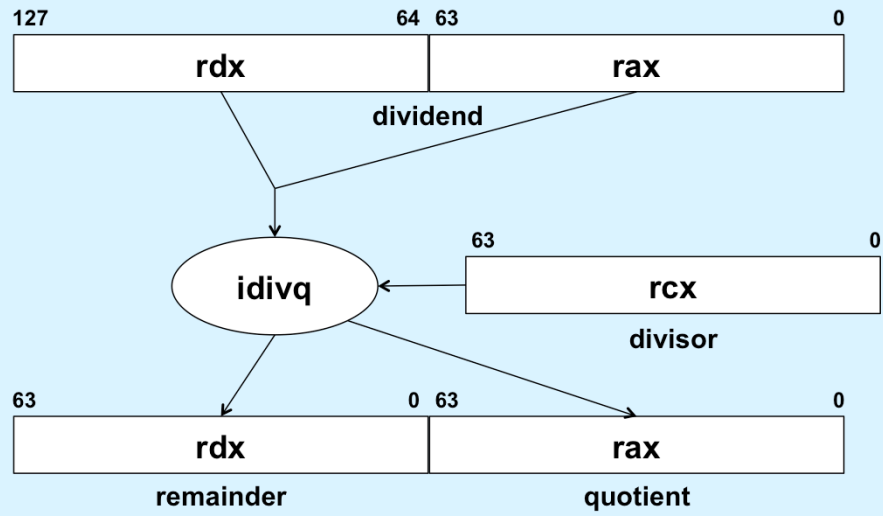
Code Blocks (Rest)

```
switch(x) {  
    . . .  
    case 2: // .L4  
        w = y/z;  
        /* Fall Through */  
    case 3: // .L9  
        w += z;  
        break;  
    . . .  
}  
  
.L4:      # x == 2  
    movq  %rsi, %rax  
    movq  %rsi, %rdx  
    sarq  $63, %rdx  
    idivq %rcx      # w = y/z  
    jmp   .L5  
.L9:      # x == 3  
    movl  $1, %eax # w = 1  
.L5:      # merge:  
    addq  %rcx, %rax # w += z  
    ret
```

Supplied by CMU, but converted to x86-64.

The code following the `.L4` label requires some explanation. The `idivq` instruction is special in that it takes a 128-bit dividend that is implicitly assumed to reside in registers `rdx` and `rax`. Its single operand specifies the divisor. The quotient is always placed in the `rax` register, and the remainder in the `rdx` register. In our example, `y`, which we want to be the dividend, is copied into both the `rax` and `rdx` registers. The `sarq` (shift arithmetic right quadword) instruction propagates the sign bit of `rdx` across the entire register, replacing its original contents. Thus, if one considers `rdx` to contain the most-significant bits of the dividend and `rax` to contain the least-significant bits, the pair of registers now contains the 128-bit version of `y`. The `idivq` instruction computes the quotient from dividing this 128-bit value by the 64-bit value contained in register `rcx` (containing `z`). The quotient is stored register `rax` (implicitly) and the remainder is stored in register `rdx` (and is ignored in our example). This illustrated in the next slide.

idivq



x86-64 Object Code

- **Setup**

- label `.L8` becomes address `0x4004e5`
- label `.L7` becomes address `0x4005c0`

Assembly code

```
switch_eg:
    . . .
    ja     .L8          # If unsigned > goto default
    jmp    *.L7(,%rdi,8) # Goto *JTab[x]
```

Disassembled object code

```
00000000004004ac <switch_eg>:
    . . .
    4004b3: 77 30                ja     4004e5 <switch_eg+0x39>
    4004b5: ff 24 fd c0 05 40 00 jmpq   *0x4005c0(,%rdi,8)
```

Supplied by CMU, but converted to x86-64.

Disassembly was accomplished using “`objdump -d`”. Note that the text enclosed in angle brackets (“<”, “>”) is essentially a comment, relating the address (4004e5) to a symbolic location (0x39 bytes after the beginning of *switch_eg*).

x86-64 Object Code (cont.)

- **Jump table**

- doesn't show up in disassembled code
- can inspect using gdb

```
gdb switch
```

```
(gdb) x/7xg 0x4005c0
```

- » examine 7 hexadecimal format “giant” words (8-bytes each)
- » use command “**help x**” to get format documentation

```
0x4005c0:      0x00000000004004e5      0x00000000004004bc
0x4005d0:      0x00000000004004c4      0x00000000004004d3
0x4005e0:      0x00000000004004e5      0x00000000004004dc
0x4005f0:      0x00000000004004dc
```

Supplied by CMU, but converted to x86-64. We assume that the `switch_eg` function was included in a program whose name is *switch*. Hence, `gdb` is invoked from the shell with the argument “switch”.

x86-64 Object Code (cont.)

- Deciphering jump table

```
0x4005c0:    0x00000000004004e5    0x00000000004004bc
0x4005d0:    0x00000000004004c4    0x00000000004004d3
0x4005e0:    0x00000000004004e5    0x00000000004004dc
0x4005f0:    0x00000000004004dc
```

Address	Value	x
0x4005c0	0x4004e5	0
0x4005c8	0x4004bc	1
0x4005d0	0x4004c4	2
0x4005d8	0x4004d3	3
0x4005e0	0x4004e5	4
0x4005e8	0x4004dc	5
0x4005f0	0x4004dc	6

Supplied by CMU, but converted to x86-64.

Disassembled Targets

```
(gdb) disassemble 0x4004bc,0x4004eb
Dump of assembler code from 0x4004bc to 0x4004eb
0x00000000004004bc <switch_eg+16>:  mov    %rsi,%rax
0x00000000004004bf <switch_eg+19>:  imul   %rdx,%rax
0x00000000004004c3 <switch_eg+23>:  retq
0x00000000004004c4 <switch_eg+24>:  mov    %rsi,%rax
0x00000000004004c7 <switch_eg+27>:  mov    %rsi,%rdx
0x00000000004004ca <switch_eg+30>:  sar    $0x3f,%rdx
0x00000000004004ce <switch_eg+34>:  idiv   %rcx
0x00000000004004d1 <switch_eg+37>:  jmp    0x4004d8 <switch_eg+44>
0x00000000004004d3 <switch_eg+39>:  mov    $0x1,%eax
0x00000000004004d8 <switch_eg+44>:  add    %rcx,%rax
0x00000000004004db <switch_eg+47>:  retq
0x00000000004004dc <switch_eg+48>:  mov    $0x1,%eax
0x00000000004004e1 <switch_eg+53>:  sub    %rdx,%rax
0x00000000004004e4 <switch_eg+56>:  retq
0x00000000004004e5 <switch_eg+57>:  mov    $0x2,%eax
0x00000000004004ea <switch_eg+62>:  retq
```

Matching Disassembled Targets

Value	x
0x4004e5	0
0x4004bc	1
0x4004c4	2
0x4004d3	3
0x4004e5	4
0x4004dc	5
0x4004dc	6

```

0x0000000004004bc:  mov    %rsi,%rax
0x0000000004004bf:  imul   %rdx,%rax
0x0000000004004c3:  retq
0x0000000004004c4:  mov    %rsi,%rax
0x0000000004004c7:  mov    %rsi,%rdx
0x0000000004004ca:  sar    $0x3f,%rdx
0x0000000004004ce:  idiv   %rcx
0x0000000004004d1:  jmp    0x4004d8
0x0000000004004d3:  mov    $0x1,%eax
0x0000000004004d8:  add    %rcx,%rax
0x0000000004004db:  retq
0x0000000004004dc:  mov    $0x1,%eax
0x0000000004004e1:  sub    %rdx,%rax
0x0000000004004e4:  retq
0x0000000004004e5:  mov    $0x2,%eax
0x0000000004004ea:  retq

```