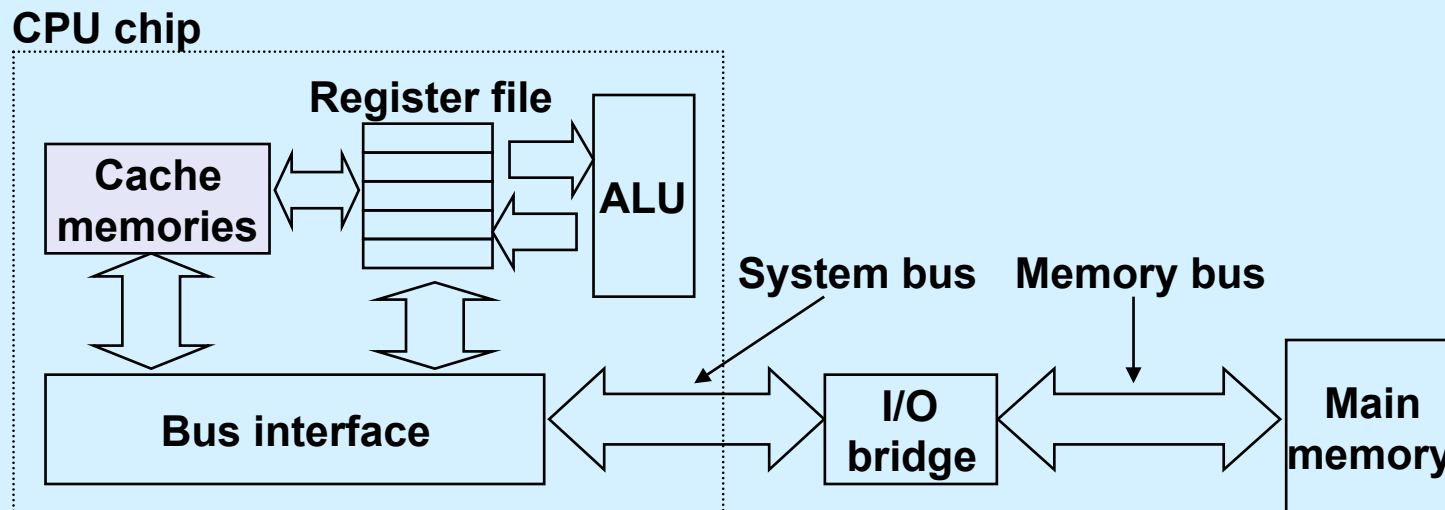


# CS 33

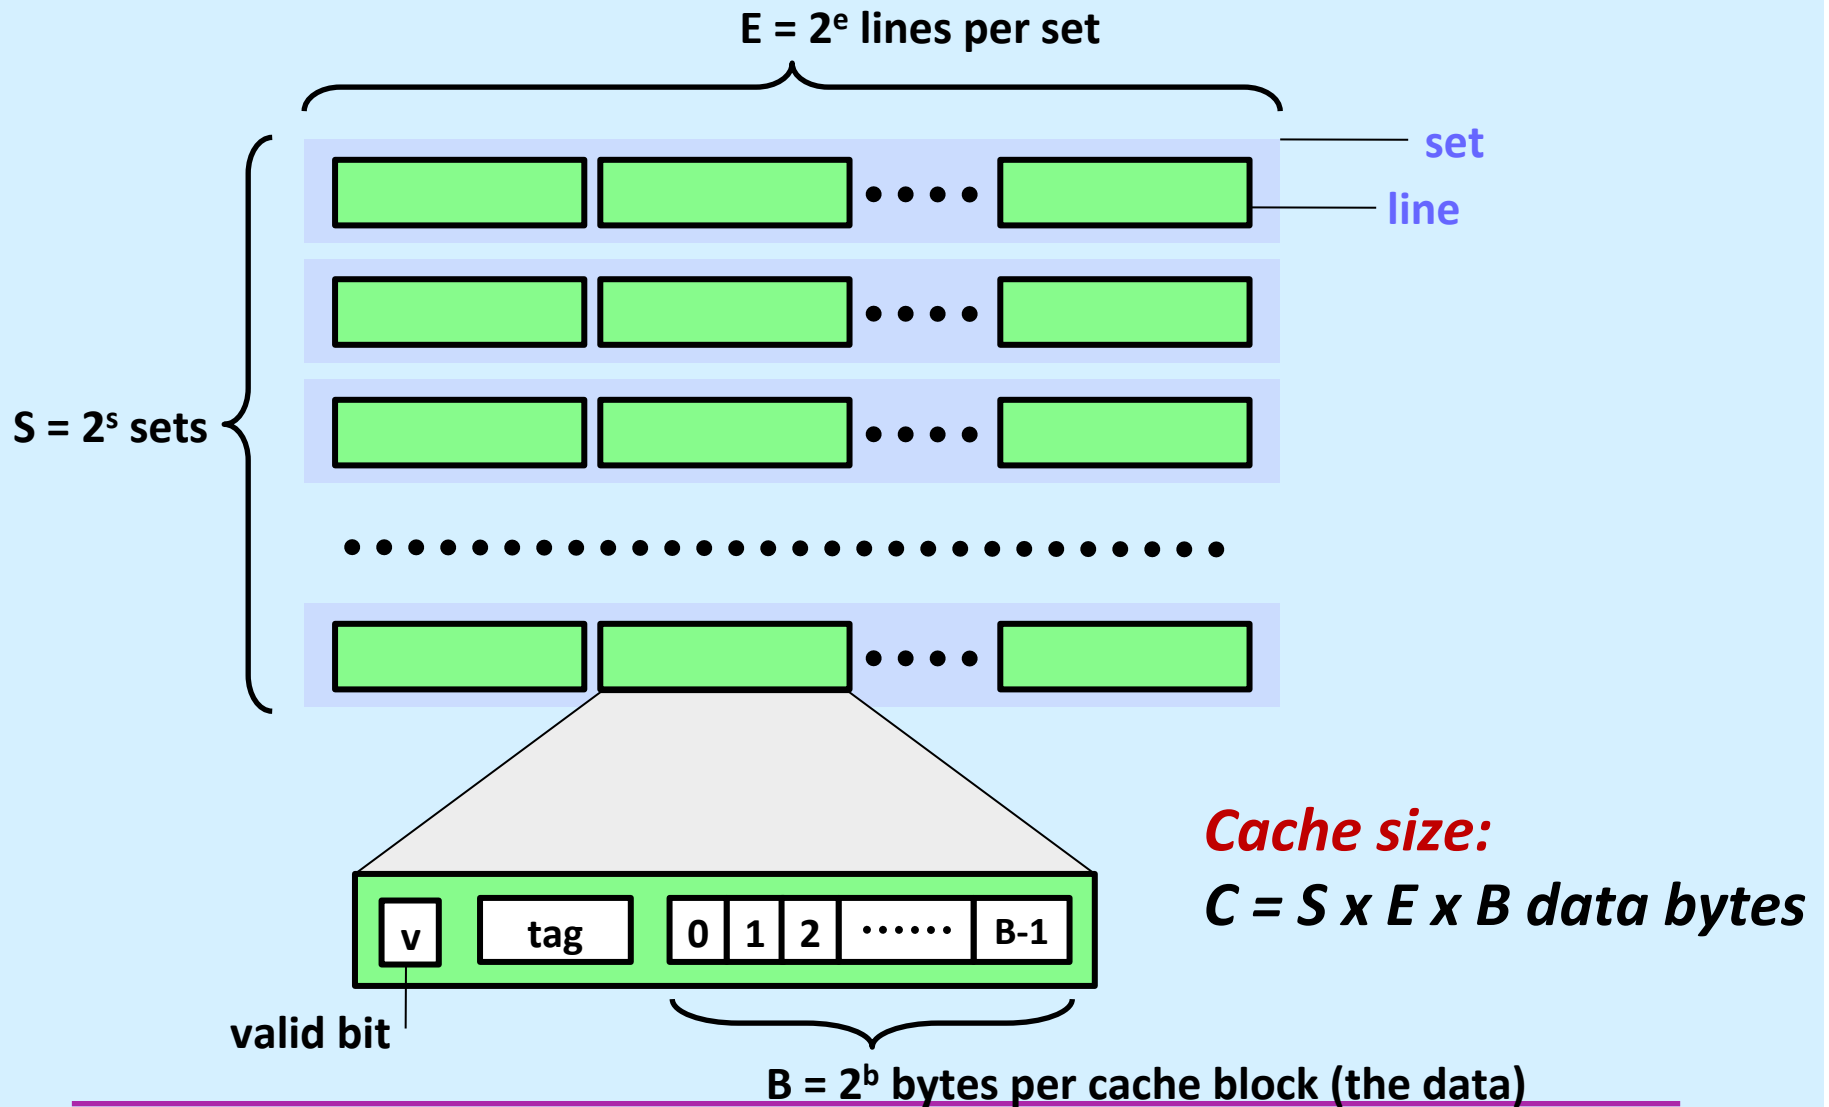
## Caches

# Cache Memories

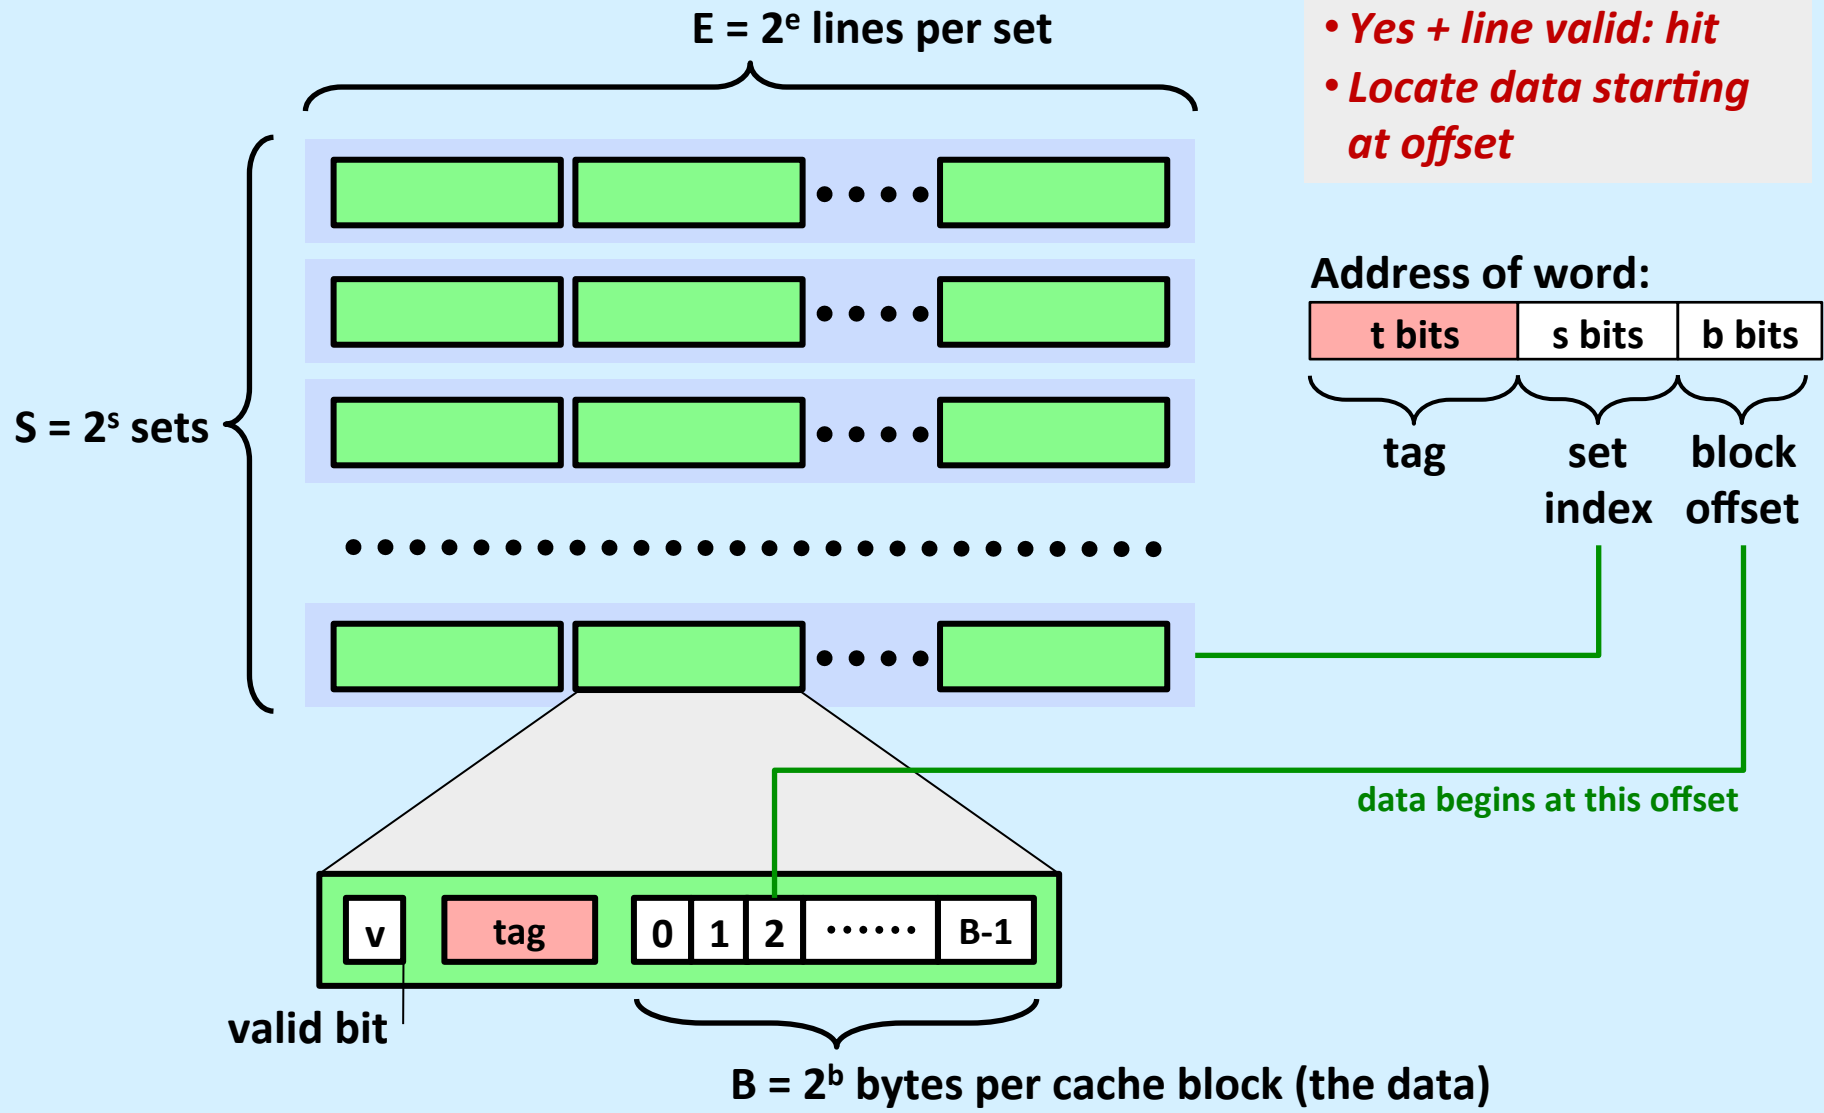
- **Cache memories** are small, fast SRAM-based memories managed automatically in hardware
  - hold frequently accessed blocks of main memory
- CPU looks first for data in caches (e.g., L1, L2, and L3), then in main memory
- Typical system structure:



# General Cache Organization (S, E, B)



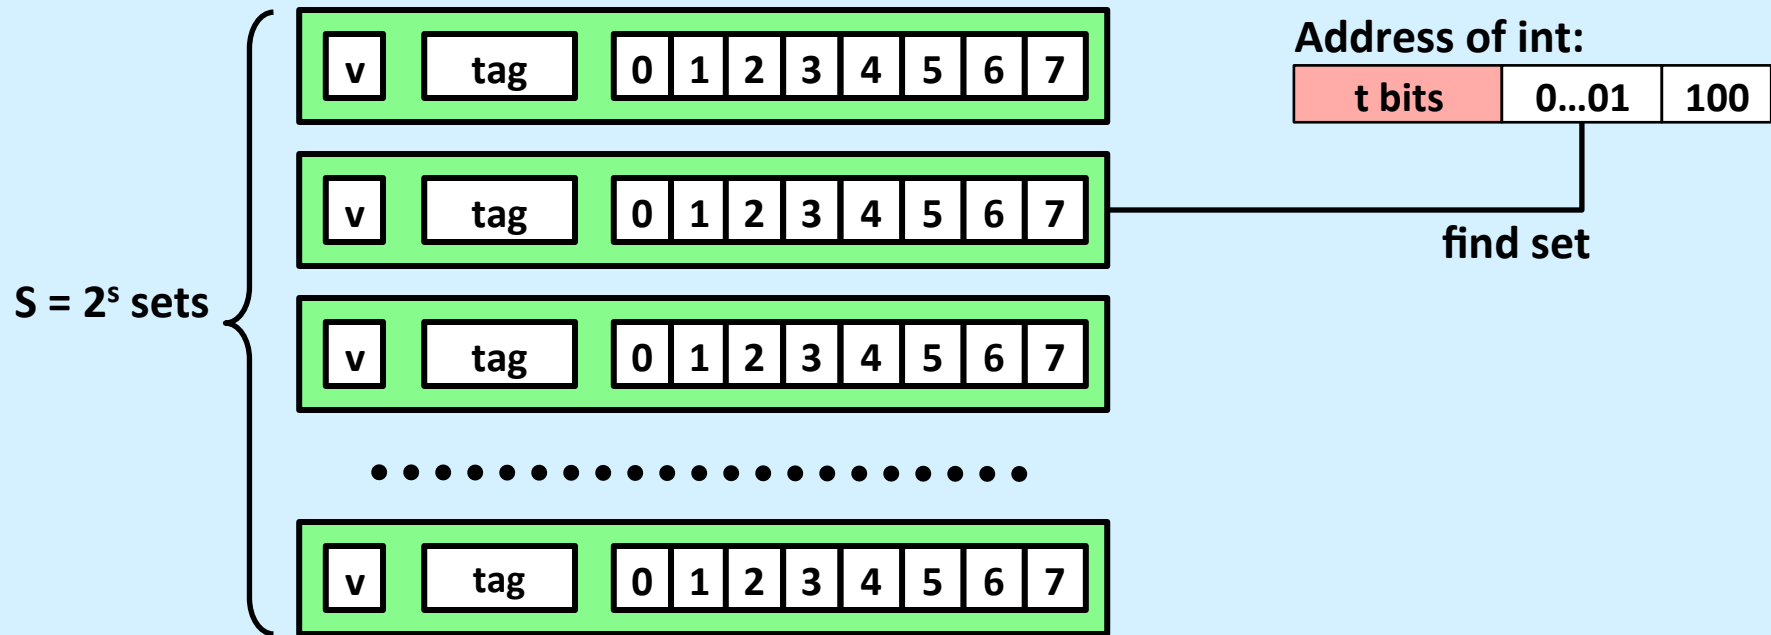
# Cache Read



# Example: Direct Mapped Cache (E = 1)

Direct mapped: one line per set


Assume: cache block size 8 bytes



# Example: Direct Mapped Cache (E = 1)

Direct mapped: one line per set

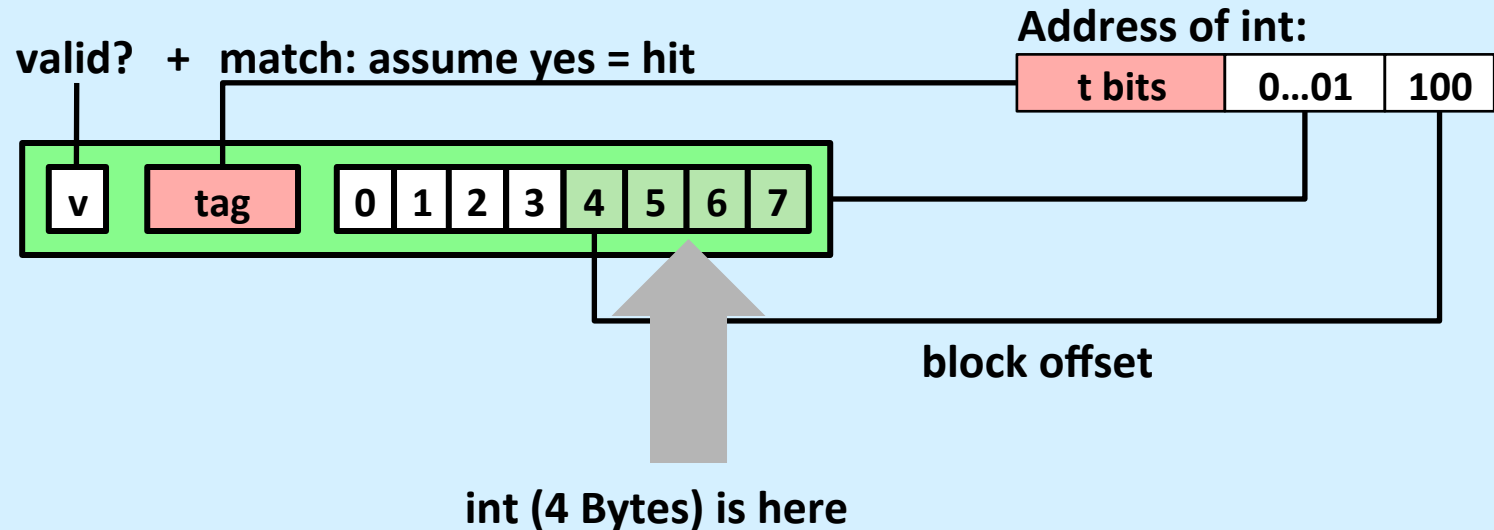
Assume: cache block size 8 bytes



# Example: Direct Mapped Cache (E = 1)

Direct mapped: one line per set

Assume: cache block size 8 bytes



**No match:** old line is evicted and replaced

# Direct-Mapped Cache Simulation

t=1	s=2	b=1
x	xx	x

M=16 byte addresses, B=2 bytes/block,  
S=4 sets, E=1 Blocks/set

Address trace (reads, one byte per read):

0	[ <u>0000</u> <sub>2</sub> ],	miss
1	[ <u>0001</u> <sub>2</sub> ],	hit
7	[ <u>0111</u> <sub>2</sub> ],	miss
8	[ <u>1000</u> <sub>2</sub> ],	miss
0	[ <u>0000</u> <sub>2</sub> ]	miss

	v	Tag	Block
Set 0	1	0	M[0-1]
Set 1			
Set 2			
Set 3	1	0	M[6-7]



# A Higher-Level Example

```
int sum_array_rows(double a[16][16])
{
    int i, j;
    double sum = 0;


    for (i = 0; i < 16; i++)
        for (j = 0; j < 16; j++)
            sum += a[i][j];
    return sum;
}
```

```
int sum_array_cols(double a[16][16])
{
    int i, j;
    double sum = 0;

    for (j = 0; j < 16; j++)
        for (i = 0; i < 16; i++)
            sum += a[i][j];
    return sum;
}
```

*Ignore the variables sum, i, j*

assume: cold (empty) cache,  
a[0][0] goes here



32 B = 4 doubles

# A Higher-Level Example

```
int sum_array_rows(double a[16][16])
{
    int i, j;
    double sum = 0;

    for (i = 0; i < 16; i++)
        for (j = 0; j < 16; j++)
            sum += a[i][j];
    return sum;
}
```

```
int sum_array_cols(double a[16][16])
{
    int i, j;
    double sum = 0;

    for (j = 0; j < 16; j++)
        for (i = 0; i < 16; i++)
            sum += a[i][j];
    return sum;
}
```

a <sub>0,0</sub>	a <sub>0,1</sub>	a <sub>0,2</sub>	a <sub>0,3</sub>
a <sub>0,4</sub>	a <sub>0,5</sub>	a <sub>0,6</sub>	a <sub>0,7</sub>
a <sub>0,8</sub>	a <sub>0,9</sub>	a <sub>0,10</sub>	a <sub>0,11</sub>
a <sub>0,12</sub>	a <sub>0,13</sub>	a <sub>0,14</sub>	a <sub>0,15</sub>
a <sub>1,0</sub>	a <sub>1,1</sub>	a <sub>1,2</sub>	a <sub>1,3</sub>
a <sub>1,4</sub>	a <sub>1,5</sub>	a <sub>1,6</sub>	a <sub>1,7</sub>
a <sub>1,8</sub>	a <sub>1,9</sub>	a <sub>1,10</sub>	a <sub>1,11</sub>
a <sub>1,12</sub>	a <sub>1,13</sub>	a <sub>1,14</sub>	a <sub>1,15</sub>



**32 B = 4 doubles**

# A Higher-Level Example

```
int sum_array_rows(double a[16][16])
{
    int i, j;
    double sum = 0;

    for (i = 0; i < 16; i++)
        for (j = 0; j < 16; j++)
            sum += a[i][j];
    return sum;
}
```

```
int sum_array_cols(double a[16][16])
{
    int i, j;
    double sum = 0;


    for (j = 0; j < 16; j++)
        for (i = 0; i < 16; i++)
            sum += a[i][j];
    return sum;
}
```



32 B = 4 doubles

# Conflict Misses: Aligned


```
double dotprod(double x[8], double y[8]) {  
    double sum = 0.0;  
    int i;  
  
    for (i=0; i<8; i++)  
        sum += x[i] * y[i];  
  
    return sum;  
}
```



32 B = 4 doubles

# Different Alignments

```
double dotprod(double x[8], double y[8]) {  
    double sum = 0.0;  
    int i;  
  
    for (i=0; i<8; i++)  
        sum += x[i] * y[i];  
  
    return sum;  
}
```

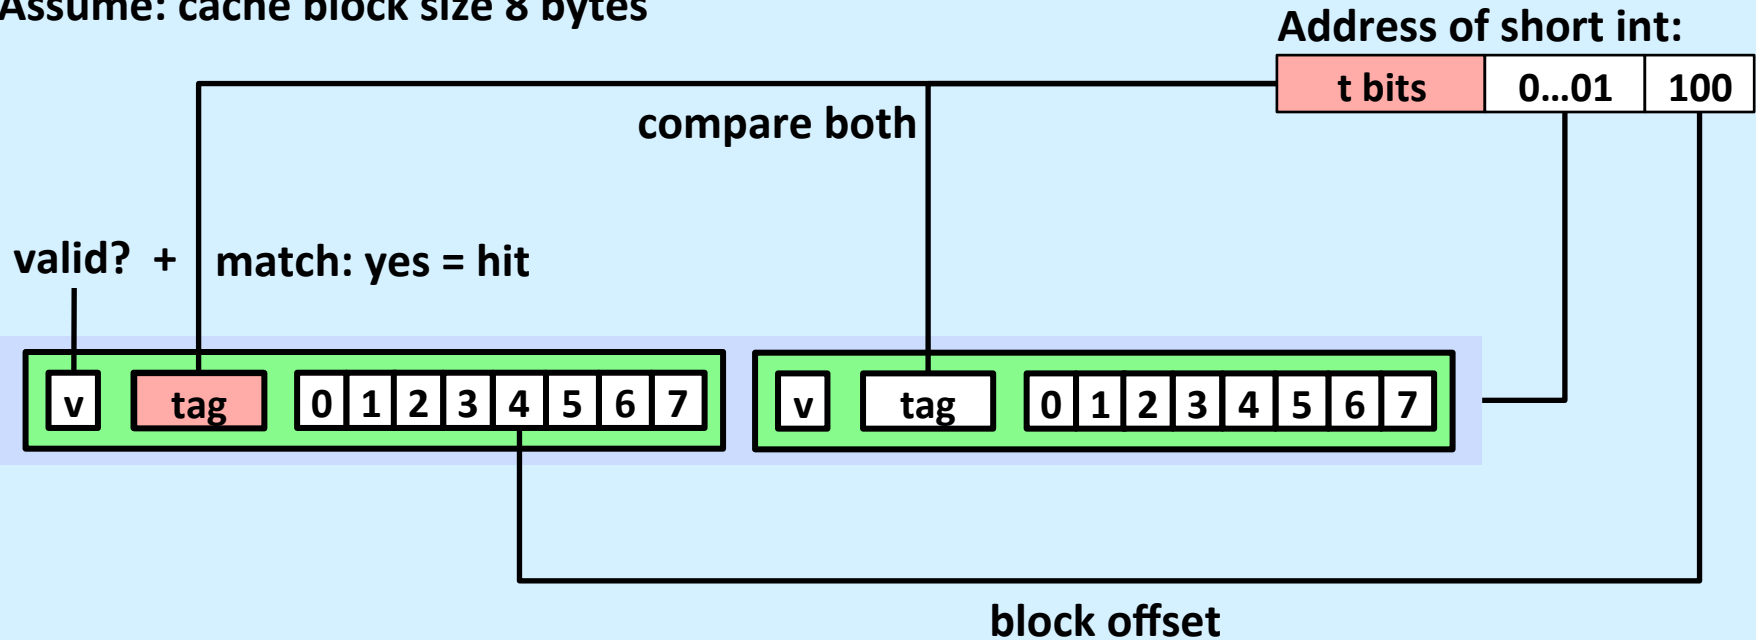


32 B = 4 doubles

# E-way Set-Associative Cache (Here: E = 2)

E = 2: two lines per set

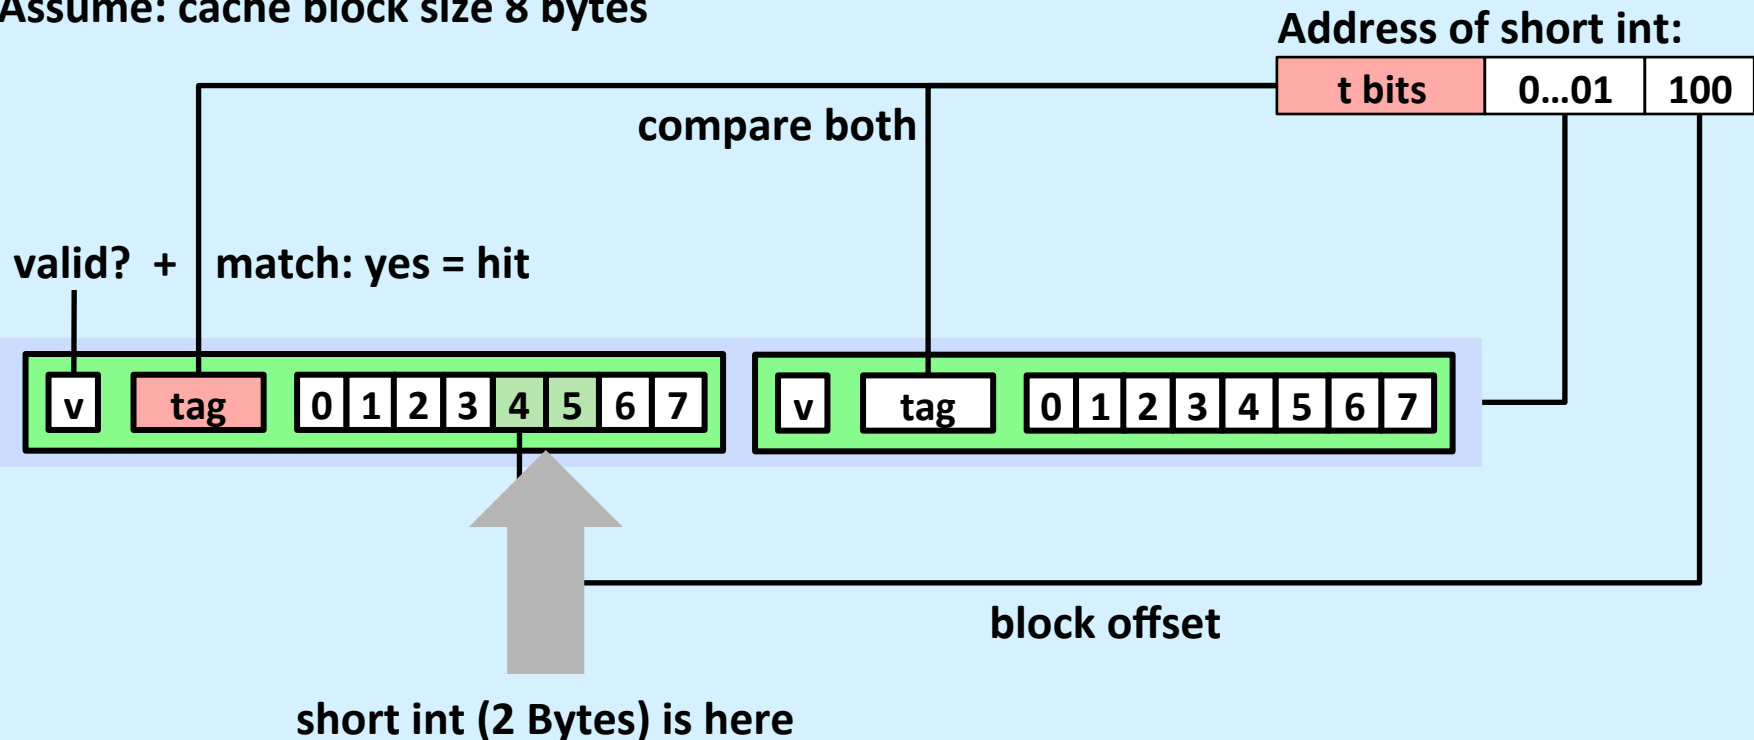
Assume: cache block size 8 bytes



# E-way Set-Associative Cache (Here: E = 2)

E = 2: two lines per set

Assume: cache block size 8 bytes



## No match:

- One line in set is selected for eviction and replacement
- Replacement policies: random, least recently used (LRU), ...

# Quiz 1

Address of int:

100	01	100
-----	----	-----

0	v	tag=0	0	0	0	0	1	1	1	1
	v	tag=2	2	2	2	2	3	3	3	3
1	v	tag=0	4	4	4	4	5	5	5	5
	v	tag=4	6	6	6	6	7	7	7	7
2	v	tag=2	8	8	8	8	9	9	9	9
	v	tag=3	a	a	a	a	b	b	b	b
3	v	tag=4	c	c	c	c	d	d	d	d
	v	tag=a	e	e	e	e	f	f	f	f

Given the address above and the cache contents as shown, what is the value of the *int* at the given address?

- a) 1111
- b) 3333
- c) 4444
- d) 7777



# 2-Way Set-Associative Cache Simulation

t=2	s=1	b=1
xx	x	x

M=16 byte addresses, B=2 bytes/block,  
S=2 sets, E=2 blocks/set

Address trace (reads, one byte per read):

0	[0000 <sub>2</sub> ],	miss
1	[0001 <sub>2</sub> ],	hit
7	[0111 <sub>2</sub> ],	miss
8	[1000 <sub>2</sub> ],	miss
0	[0000 <sub>2</sub> ]	hit

	v	Tag	Block
Set 0	1	00	M[0-1]
	1	10	M[8-9]
Set 1	1	01	M[6-7]
	0		

# A Higher-Level Example

```
int sum_array_rows(double a[16][16])
{
    int i, j;
    double sum = 0;


    for (i = 0; i < 16; i++)
        for (j = 0; j < 16; j++)
            sum += a[i][j];
    return sum;
}
```

```
int sum_array_rows(double a[16][16])
{
    int i, j;
    double sum = 0;

    for (j = 0; j < 16; j++)
        for (i = 0; i < 16; i++)
            sum += a[i][j];
    return sum;
}
```

*Ignore the variables sum, i, j*

assume: cold (empty) cache,  
a[0][0] goes here



32 B = 4 doubles

# A Higher-Level Example


*Ignore the variables sum, i, j*

```
int sum_array_rows(double a[16][16])
{
    int i, j;
    double sum = 0;

    for (i = 0; i < 16; i++)
        for (j = 0; j < 16; j++)
            sum += a[i][j];
    return sum;
}
```

```
int sum_array_cols(double a[16][16])
{
    int i, j;
    double sum = 0;

    for (j = 0; j < 16; j++)
        for (i = 0; i < 16; i++)
            sum += a[i][j];
    return sum;
}
```



**32 B = 4 doubles**

# A Higher-Level Example


*Ignore the variables sum, i, j*

```
int sum_array_rows(double a[16][16])
{
    int i, j;
    double sum = 0;

    for (i = 0; i < 16; i++)
        for (j = 0; j < 16; j++)
            sum += a[i][j];
    return sum;
}
```


```
int sum_array_cols(double a[16][16])
{
    int i, j;
    double sum = 0;

    for (j = 0; j < 16; j++)
        for (i = 0; i < 16; i++)
            sum += a[i][j];
    return sum;
}
```



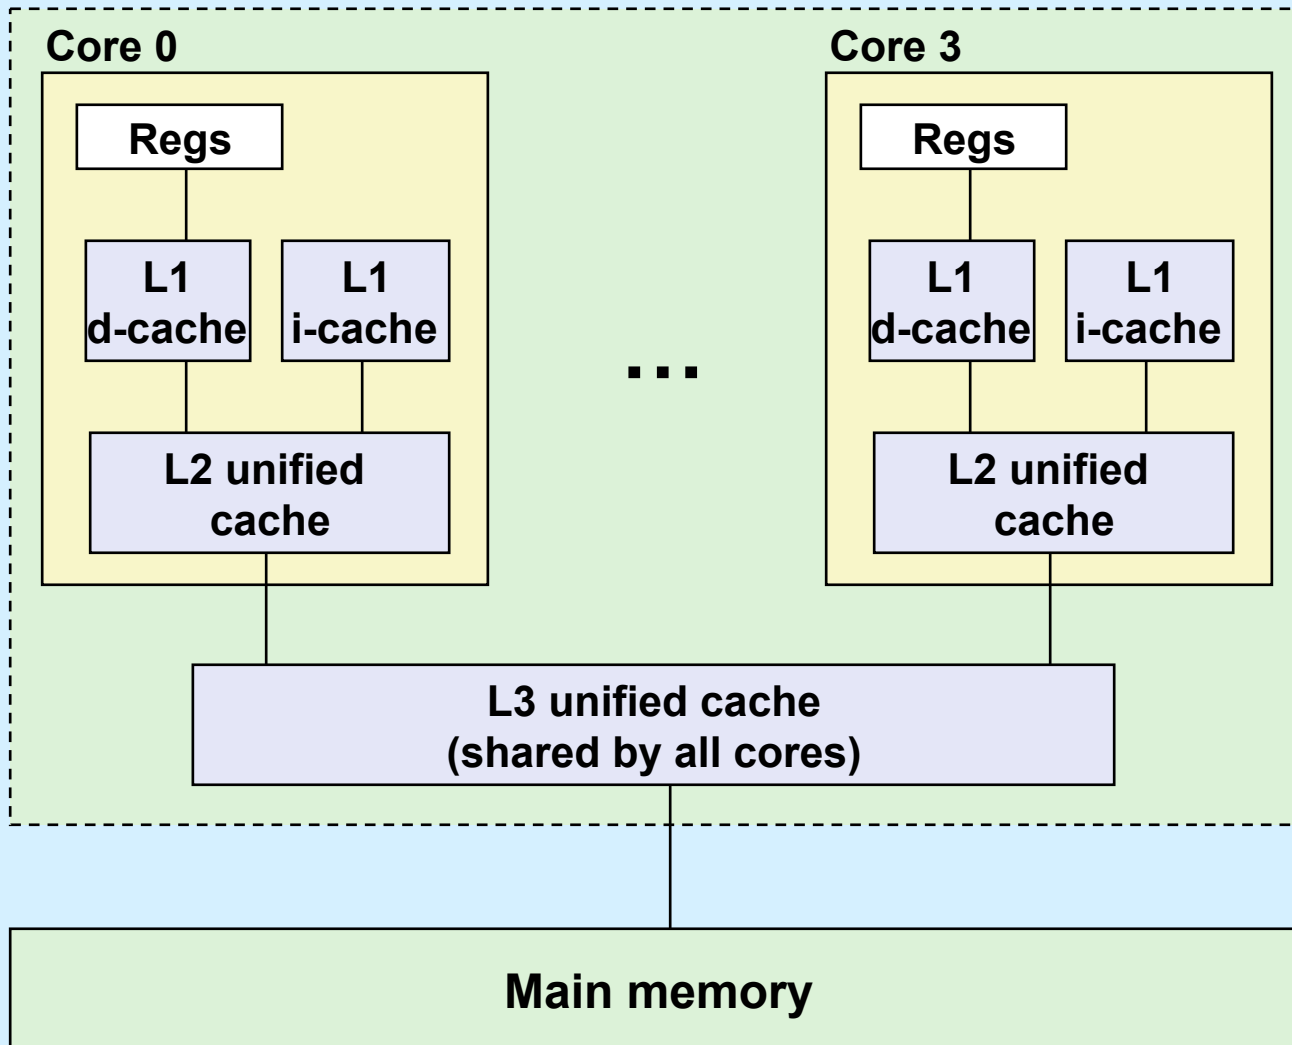
# Conflict Misses

```
double dotprod(double x[8], double y[8]) {  
    double sum = 0.0;  
    int i;  
  
    for (i=0; i<8; i++)  
        sum += x[i] * y[i];  
  
    return sum;  
}
```



# Intel Core i7 Cache Hierarchy

## Processor package



**L1 i-cache and d-cache:**  
32 KB, 8-way,  
Access: 4 cycles

**L2 unified cache:**  
256 KB, 8-way,  
Access: 11 cycles

**L3 unified cache:**  
8 MB, 16-way,  
Access: 30-40 cycles

**Block size:** 64 bytes for  
all caches

# What About Writes?

- Multiple copies of data exist:
    - L1, L2, main memory, disk
  - What to do on a write-hit?
    - **write-through** (write immediately to memory)
    - **write-back** (defer write to memory until replacement of line)
      - » need a dirty bit (line different from memory or not)
  - What to do on a write-miss?
    - **write-allocate** (load into cache, update line in cache)
      - » good if more writes to the location follow
    - **no-write-allocate** (writes immediately to memory)
  - Typical
    - write-through + no-write-allocate
    - write-back + write-allocate
-

# Cache Performance Metrics

- **Miss rate**
    - fraction of memory references not found in cache (misses / accesses)  
= 1 – hit rate
    - typical numbers (in percentages):
      - » 3-10% for L1
      - » can be quite small (e.g., < 1%) for L2, depending on size, etc.
  - **Hit time**
    - time to deliver a line in the cache to the processor
      - » includes time to determine whether the line is in the cache
    - typical numbers:
      - » 1-2 clock cycles for L1
      - » 5-20 clock cycles for L2
  - **Miss penalty**
    - additional time required because of a miss
      - » typically 50-200 cycles for main memory (trend: increasing!)
-



# Let's Think About Those Numbers

- Huge difference between a hit and a miss
  - could be 100x, if just L1 and main memory
- Would you believe 99% hit rate is twice as good as 97%?
  - consider:
    - cache hit time of 1 cycle
    - miss penalty of 100 cycles
  - average access time:
    - 97% hits:  $.97 * 1 \text{ cycle} + 0.03 * 100 \text{ cycles} \approx 4 \text{ cycles}$
    - 99% hits:  $.99 * 1 \text{ cycle} + 0.01 * 100 \text{ cycles} \approx 2 \text{ cycles}$
- This is why “miss rate” is used instead of “hit rate”


# Writing Cache-Friendly Code

- **Make the common case go fast**
  - focus on the inner loops of the core functions
- **Minimize the misses in the inner loops**
  - repeated references to variables are good (**temporal locality**)
  - stride-1 reference patterns are good (**spatial locality**)

**Key idea: our qualitative notion of locality is quantified through our understanding of cache memories**

# Miss-Rate Analysis for Matrix Multiply

- **Assume:**
  - Block size = 32B (big enough for four 64-bit words)
  - matrix dimension (N) is very large
    - » approximate  $1/N$  as 0.0
  - cache is not big enough to hold multiple rows
- **Analysis method:**
  - look at access pattern of inner loop



# Matrix Multiplication Example

- **Description:**
  - multiply  $N \times N$  matrices
  - $O(N^3)$  total operations
  - $N$  reads per source element
  - $N$  values summed per destination
    - » but may be able to hold in register

```
/* ijk */  
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum;  
    }  
}
```

*Variable sum  
held in register*


# Layout of C Arrays in Memory (review)

- **C arrays allocated in row-major order**
  - each row in contiguous memory locations
- **Stepping through columns in one row:**
  - `for (i = 0; i < N; i++)`  
    `sum += a[0][i];`
  - **accesses successive elements**
  - **if block size (B) > 4 bytes, exploit spatial locality**
    - » compulsory miss rate = 4 bytes / B
- **Stepping through rows in one column:**
  - `for (i = 0; i < n; i++)`  
    `sum += a[i][0];`
  - **accesses distant elements**
  - **no spatial locality!**
    - » compulsory miss rate = 1 (i.e. 100%)

# Matrix Multiplication (ijk)

```
/* ijk */
for (i=0; i<n; i++) {
    for (j=0; j<n; j++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum;
    }
}
```

Inner loop:

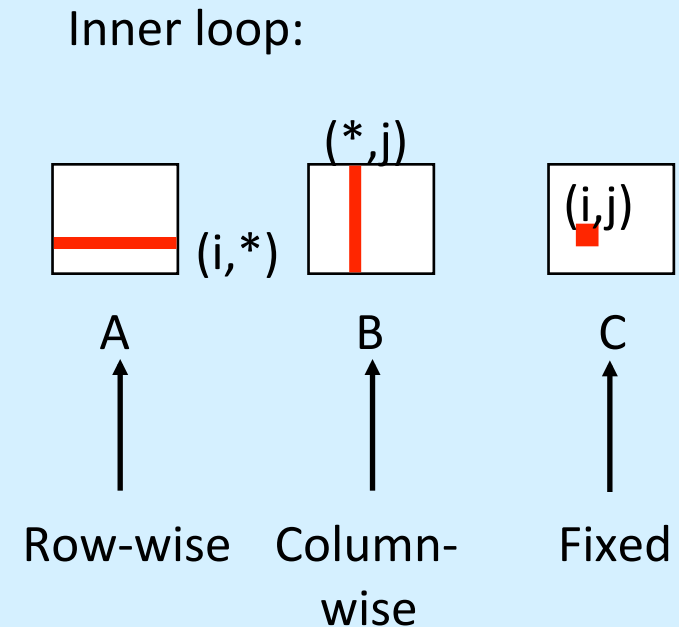


Misses per inner loop iteration:

<u>A</u>	<u>B</u>	<u>C</u>
0.25	1.0	0.0

# Matrix Multiplication (jik)

```
/* jik */
for (j=0; j<n; j++) {
    for (i=0; i<n; i++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum
    }
}
```




Misses per inner loop iteration:

<u>A</u>	<u>B</u>	<u>C</u>
0.25	1.0	0.0

# Matrix Multiplication (kij)

```
/* kij */  
for (k=0; k<n; k++) {  
    for (i=0; i<n; i++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```

Inner loop:



Misses per inner loop iteration:

A  
0.0


B  
0.25

C  
0.25



# Matrix Multiplication (ikj)

```
/* ikj */  
for (i=0; i<n; i++) {  
    for (k=0; k<n; k++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```



Misses per inner loop iteration:

A  
0.0

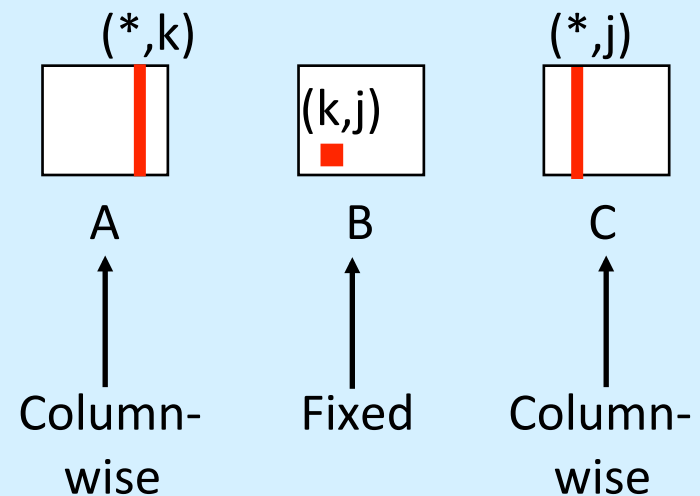
B  
0.25

C  
0.25

# Matrix Multiplication (jki)

```
/* jki */  
for (j=0; j<n; j++) {  
    for (k=0; k<n; k++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```

Inner loop:



Misses per inner loop iteration:


A  
1.0

B  
0.0

C  
1.0

# Matrix Multiplication (kji)

```
/* kji */  
for (k=0; k<n; k++) {  
    for (j=0; j<n; j++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```



Misses per inner loop iteration:

A  
1.0

B  
0.0

C  
1.0

# Summary of Matrix Multiplication

```
for (i=0; i<n; i++)  
  for (j=0; j<n; j++) {  
    sum = 0.0;  
    for (k=0; k<n; k++)  
      sum += a[i][k] * b[k][j];  
    c[i][j] = sum;  
  }
```

**ijk (& jik):**

- 2 loads, 0 stores
- misses/iter = **1.25**

```
for (k=0; k<n; k++)  
  for (i=0; i<n; i++) {  
    r = a[i][k];  
    for (j=0; j<n; j++)  
      c[i][j] += r * b[k][j];  
  }
```

**kij (& ikj):**


- 2 loads, 1 store
- misses/iter = **0.5**

```
for (j=0; j<n; j++)  
  for (k=0; k<n; k++) {  
    r = b[k][j];  
    for (i=0; i<n; i++)  
      c[i][j] += a[i][k] * r;  
  }
```

**jki (& kji):**


- 2 loads, 1 store
- misses/iter = **2.0**

# Core i7 Matrix Multiply Performance



# Matrix Multiplication: More Analysis

```
/* Multiply n x n matrices a and b */  
void mmm(double *a, double *b, double *c, int n) {  
    int i, j, k;  
    for (i = 0; i < n; i++)  
        for (j = 0; j < n; j++)  
            for (k = 0; k < n; k++)  
                c[i*n+j] += a[i*n + k]*b[k*n + j];  
}
```




# Cache-Miss Analysis


- Assume:
  - matrix elements are doubles
  - cache block = 8 doubles
  - cache size  $C \ll n$  (much smaller than  $n$ )

- First iteration:

- $n/8 + n = 9n/8$  misses



- afterwards **in cache:**  
(schematic)




# Cache-Miss Analysis

- Assume:
  - matrix elements are doubles
  - cache block = 8 doubles
  - cache size  $C \ll n$  (much smaller than  $n$ )

- Second iteration:

- again:  
 $n/8 + n = 9n/8$  misses




- Total misses:
  - $9n/8 * n^2 = (9/8) * n^3$




# Blocked Matrix Multiplication

```
/* Multiply n x n matrices a and b */
void mmm(double *a, double *b, double *c, int n) {
    int i, j, k;
    for (i = 0; i < n; i+=B)
        for (j = 0; j < n; j+=B)
            for (k = 0; k < n; k+=B)
                /* B x B mini matrix multiplications */
                for (i1 = i; i1 < i+B; i++)
                    for (j1 = j; j1 < j+B; j++)
                        for (k1 = k; k1 < k+B; k++)
                            c[i1*n+j1] += a[i1*n + k1]*b[k1*n + j1];
}
```




# Cache-Miss Analysis

- Assume:
  - cache block = 8 doubles
  - cache size  $C \ll n$  (much smaller than  $n$ )
  - three blocks  fit into cache:  $3B^2 < C$


- First (block) iteration:

- $B^2/8$  misses for each block
- $2n/B * B^2/8 = nB/4$   
(omitting matrix  $c$ )




- afterwards in cache  
(schematic)

# Cache-Miss Analysis

- Assume:
  - cache block = 8 doubles
  - cache size  $C \ll n$  (much smaller than  $n$ )
  - three blocks  fit into cache:  $3B^2 < C$

- Second (block) iteration:

- same as first iteration
- $2n/B * B^2/8 = nB/4$



- Total misses:

- $nB/4 * (n/B)^2 = n^3/(4B)$

# Summary

- **No blocking:**  $(9/8) * n^3$
- **Blocking:**  $1/(4B) * n^3$
- **Suggest largest possible block size B, but limit  $3B^2 < C$ !**
- **Reason for dramatic difference:**
  - **matrix multiplication has inherent temporal locality:**
    - » **input data:**  $3n^2$ , **computation**  $2n^3$
    - » **every array element used  $O(n)$  times!**
  - **but program has to be written properly**

# Quiz 2

**What is the smallest value of  $B$  (in 8-byte doubles) for which the cache-miss analysis works?**

- a) 1
- b) 2
- c) 4
- d) 8

# Concluding Observations

- **Programmer can optimize for cache performance**
  - how data structures are organized
  - how data are accessed
    - » nested loop structure
    - » blocking is a general technique
- **All systems favor “cache-friendly code”**
  - getting absolute optimum performance is very platform specific
    - » cache sizes, line sizes, associativities, etc.
  - can get most of the advantage with generic code
    - » keep working set reasonably small (temporal locality)
    - » use small strides (spatial locality)