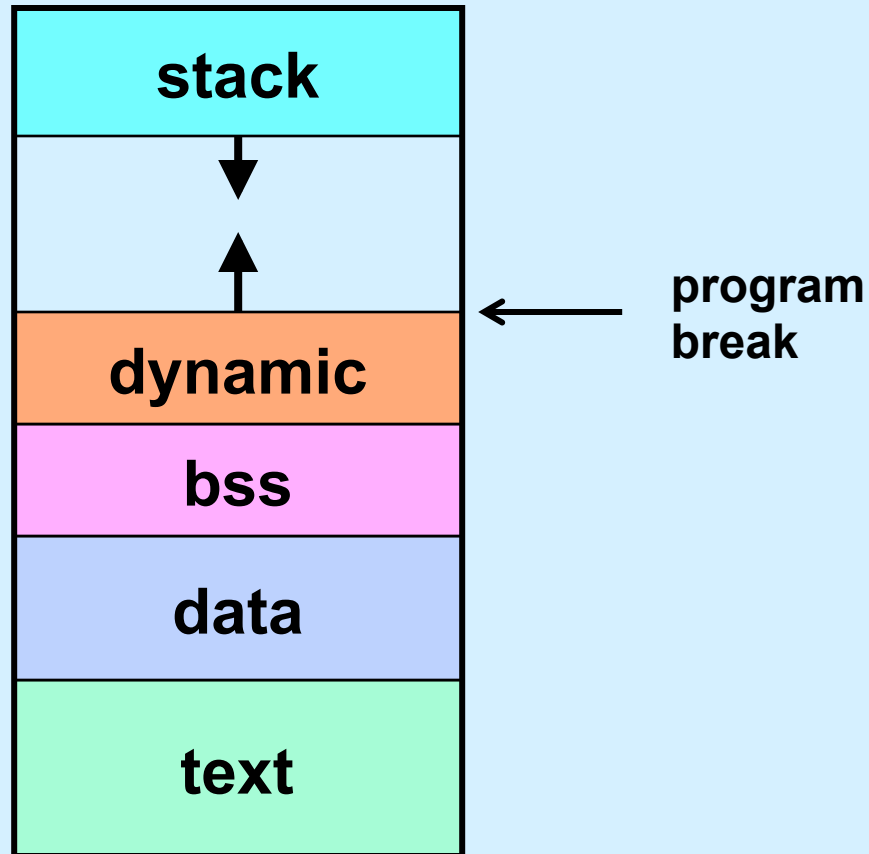


CS 33

Storage Allocation

The Unix Address Space



sbrk System Call

```
void *sbrk(intptr_t increment)
```

- moves the program break by an amount equal to *increment*
- returns the previous program break
- *intptr_t* is typedef'd to be a *long*

Managing Dynamic Storage


- **Strategy**
 - get a “chunk” of memory from the OS using *sbrk*
 - » create pool of available storage, aka the “heap”
 - *malloc*, *calloc*, *realloc*, and *free* use this storage if possible
 - » they manage the heap
 - if not possible, get more storage from OS
 - » heap is made larger (by calling *sbrk*)
- **Important note:**
 - when process terminates, all storage is given back to the system
 - » all memory-related sins are forgotten!

Dynamic Memory Allocation

- Allocator maintains heap as collection of variable sized **blocks**, which are either **allocated** or **free**
- Types of allocators
 - **explicit allocator**: application allocates and frees space
 - » e.g., malloc and free in C
 - **implicit allocator**: application allocates, but does not free space
 - » e.g. garbage collection in Java, ML, and Racket

Assumptions Made in This Lecture

- Memory is word addressed (each word can hold a pointer)




Allocation Example


`p1 = malloc(4)`




`p2 = malloc(5)`




`p3 = malloc(6)`



`free(p2)`



`p4 = malloc(2)`




Constraints

- **Applications**
 - can issue arbitrary sequence of `malloc` and `free` requests
 - `free` request must be to a `malloc`'d block
- **Allocators**
 - can't control number or size of allocated blocks
 - must respond immediately to `malloc` requests
 - » i.e., can't reorder or buffer requests
 - must allocate blocks from free memory
 - » i.e., can only place allocated blocks in free memory
 - must align blocks so they satisfy all alignment requirements
 - » 8-byte alignment for GNU `malloc` (`libc malloc`) on Linux boxes
 - can manipulate and modify only free memory
 - can't move the allocated blocks once they are `malloc`'d
 - » i.e., compaction is not allowed

Internal Fragmentation

- For a given block, *internal fragmentation* occurs if payload is smaller than block size




- Caused by
 - overhead of maintaining heap data structures
 - padding for alignment purposes
 - explicit policy decisions
(e.g., to return a big block to satisfy a small request)
- Depends only on the pattern of *previous* requests
 - thus, easy to measure


External Fragmentation

- Occurs when there is enough aggregate heap memory, but no single free block is large enough


`p1 = malloc(4)`




`p2 = malloc(5)`



`p3 = malloc(6)`



`free(p2)`



`p4 = malloc(6)`

Oops! (what would happen now?)


- Depends on the pattern of future requests
 - thus, difficult to measure

Implementation Issues

- How do we know how much memory to free given just a pointer?
- How do we keep track of the free blocks?
- What do we do with the extra space when allocating a structure that is smaller than the free block it is placed in?
- How do we pick a block to use for allocation — many might fit?
- How do we reinsert freed block?


Knowing How Much to Free

- Standard method
 - keep the length of a block in the word preceding the block
 - » this word is often called the *header field* or *header*
 - requires an extra word for every allocated block




Keeping Track of Free Blocks

- Method 1: *Implicit list* using length—links all blocks



- Method 2: *Explicit list* among the free blocks using pointers




- Method 3: *Segregated free list*
 - different free lists for different size classes
- Method 4: *Blocks sorted by size*
 - can use a balanced tree (e.g. red-black tree) with pointers within each free block, and the length used as a key

Method 1: Implicit List

- For each block we need both size and allocation status
 - could store this information in two words: wasteful!
- Standard trick
 - if blocks are aligned, some low-order address bits are always 0
 - instead of storing an always-0 bit, use it as a allocated/free flag
 - when reading size word, mask out this bit

*Format of
allocated and
free blocks*




a = 1: Allocated block
a = 0: Free block

Size: block size

Payload: application data
(allocated blocks only)

Detailed Implicit Free-List Example



Double-word
aligned

Allocated blocks: shaded
Free blocks: unshaded
Headers: labeled with size in bytes/allocated bit

Implicit List: Finding a Free Block

- **First fit:**

- search list from beginning, choose **first** free block that fits:

```
p = start;
while ((p < end) &&           // not past end
        ((*p & 1) ||          // already allocated
        (*p <= len)))         // too small
    p = p + (*p & -2);         // goto next block (word addressed)
```

- can take linear time in total number of blocks (allocated and free)
- in practice it can cause “splinters” at beginning of list


- **Next fit:**

- like first fit, but search list starting where previous search finished
- should often be faster than first fit: avoids re-scanning unhelpful blocks
- some research suggests that fragmentation is worse

- **Best fit:**

- search the list, choose the **best** free block: fits, with fewest bytes left over
- keeps fragments small—usually helps fragmentation
- will typically run slower than first fit


Quiz 1



We have two free blocks of memory, of sizes 1300 and 1200 (appearing in that order). There are three successive requests to *malloc* for allocations of 1000, 1100, and 250 bytes. Which approach does best? (Hint: one of the two fails the last request.)

- a) first fit**
- b) best fit**


Allocation



1300

1200


First Fit



300

1000 bytes


1200



300

1100 bytes

100



50

250 bytes


100

Best Fit




1300

200



200


200




Stuck!

Implicit List: Allocating in Free Block

- Allocating in a free block: *splitting*
 - since allocated space might be smaller than free space, we might want to split the block



`addblock(p, 4)`



```
void addblock(ptr p, int len) {  
    int newsize = ((len + 1) >> 1) << 1; // round up to even  
    int oldsize = *p & -2;                // mask out low bit  
    *p = newsize | 1;                     // set new length  
    if (newsize < oldsize)  
        *(p+newsize) = oldsize - newsize; // set length in remaining  
}
```

// part of block


Implicit List: Freeing a Block

- Simplest implementation:

- need only clear the “allocated” flag

```
void free_block(ptr p) { *p = *p & -2 }
```

- but can lead to “false fragmentation”




`malloc(5)` ***Oops!***

There is enough free space, but the allocator won't be able to find it

Implicit List: Coalescing

- Join (**coalesce**) with next/previous blocks, if they are free
 - coalescing with next block




```
void free_block(ptr p) {  
    *p = *p & -2;           // clear allocated flag  
    next = p + *p;          // find next block  
    if ((*next & 1) == 0)   // add to this block if  
        *p = *p + *next;    // not allocated  
}
```

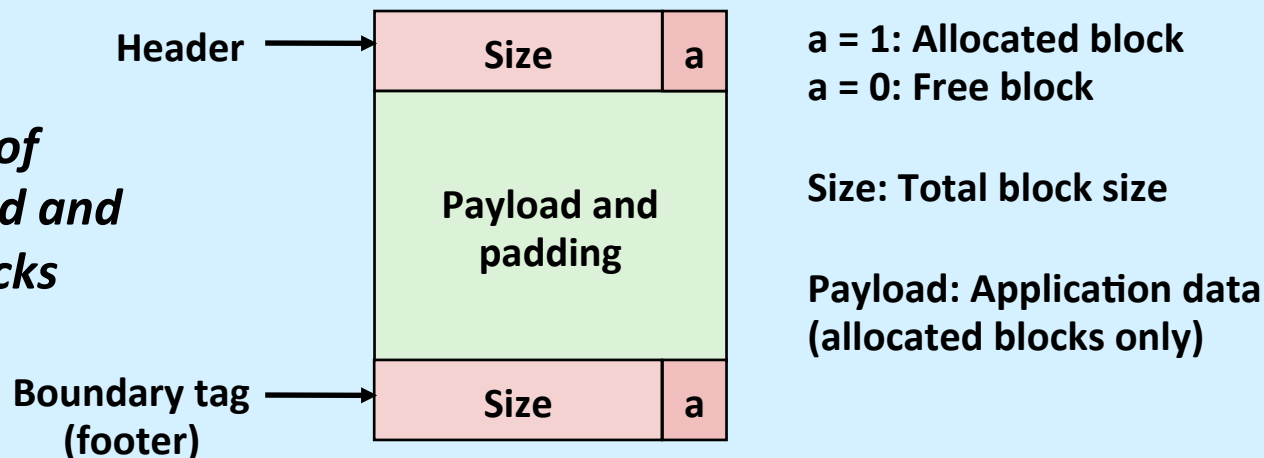
– but how do we coalesce with *previous* block?

Implicit List: Bidirectional Coalescing

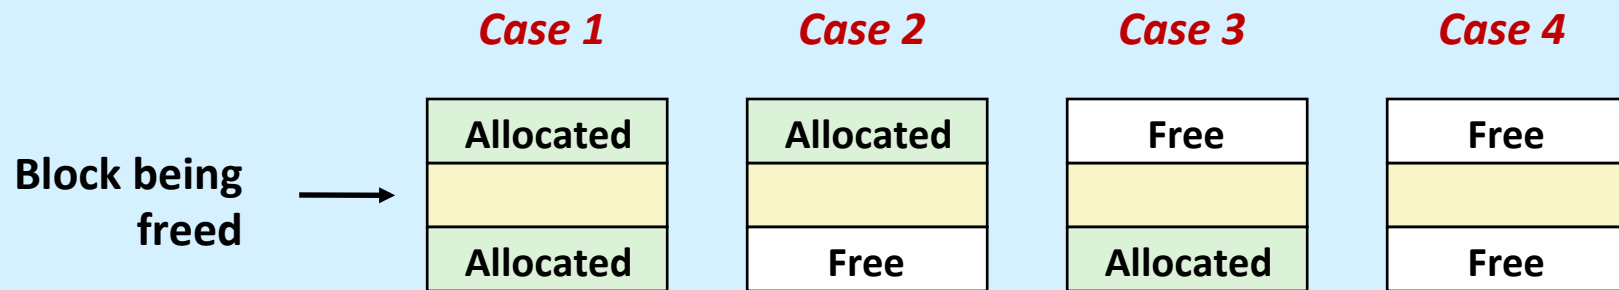
- **Boundary tags** [Knuth73]
 - replicate size/allocated word at “bottom” (end) of free blocks
 - allows us to traverse the “list” backwards, but requires extra space
 - important and general technique!




*Format of
allocated and
free blocks*



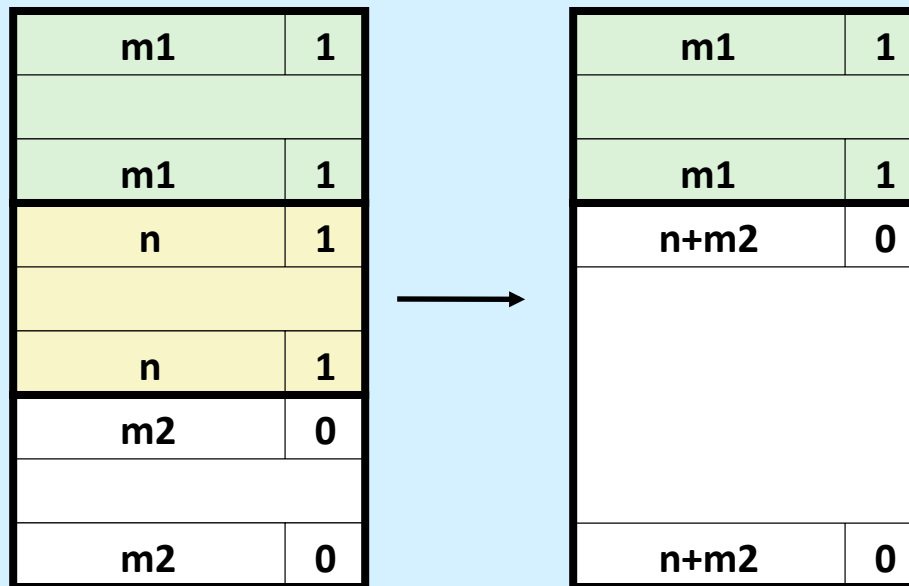
Constant Time Coalescing



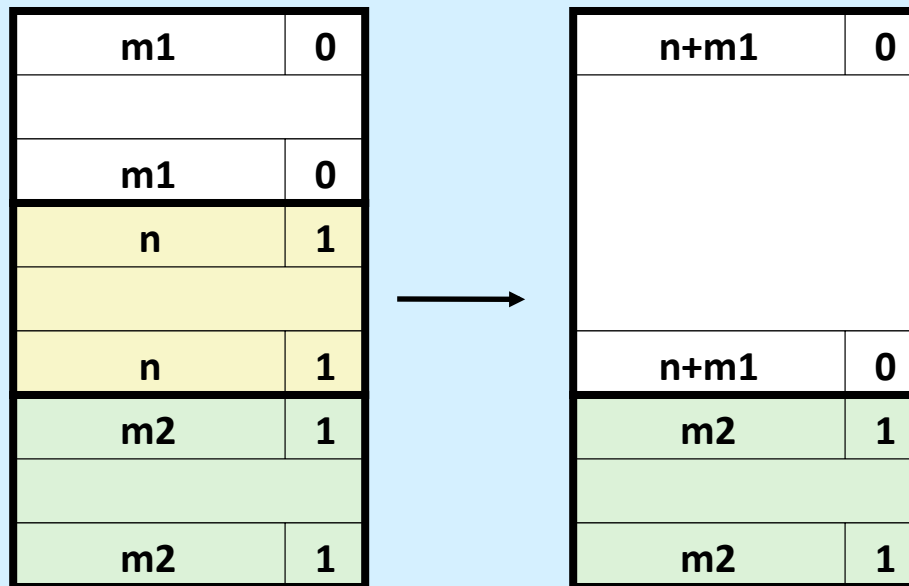
Constant Time Coalescing (Case 1)




Constant Time Coalescing (Case 2)



Constant Time Coalescing (Case 3)



Constant Time Coalescing (Case 4)



Summary of Key Allocator Policies


- Placement policy:
 - first-fit, next-fit, best-fit, etc.
 - trades off lower throughput for less fragmentation
 - **interesting observation**: segregated free lists approximate a best-fit placement policy without having to search entire free list
- Splitting policy:
 - when do we go ahead and split free blocks?
 - how much internal fragmentation are we willing to tolerate?
- Coalescing policy:
 - **immediate coalescing**: coalesce each time `free` is called
 - **deferred coalescing**: try to improve performance of `free` by deferring coalescing until needed. Examples:
 - » coalesce as you scan the free list for `malloc`
 - » coalesce when the amount of external fragmentation reaches some threshold

Implicit Lists: Summary


- **Implementation: very simple**
- **Allocate cost:**
 - linear time worst case
- **Free cost:**
 - constant time worst case
 - even with coalescing
- **Memory usage:**
 - will depend on placement policy
 - first-fit, next-fit or best-fit
- **Not used in practice for `malloc/free` because of linear-time allocation**
 - used in many special purpose applications
- **However, the concepts of splitting and boundary tag coalescing are general to *all* allocators**

Keeping Track of Free Blocks

- Method 1: *implicit free list* using length—links all blocks




- Method 2: *explicit free list* among the free blocks using pointers



- Method 3: *segregated free list*
 - different free lists for different size classes
- Method 4: *blocks sorted by size*
 - can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

Explicit Free Lists

Allocated (as before)




Free




- Maintain list(s) of **free** blocks, not **all** blocks
 - the “next” free block could be anywhere
 - » so we need to store forward/back pointers, not just sizes
 - » luckily we track only free blocks, so we can use payload area
 - still need boundary tags for coalescing

Explicit Free Lists

- Logically:




- Physically: blocks can be in any order



Allocating From Explicit Free Lists


conceptual graphic

Before



After

(with splitting)




Freeing With Explicit Free Lists

- **Insertion policy:** where in the free list do you put a newly freed block?
 - LIFO (last-in-first-out) policy
 - » insert freed block at the beginning of the free list
 - » **pro:** simple and constant time
 - » **con:** studies suggest fragmentation is worse than address ordered
 - address-ordered policy
 - » Insert freed blocks so that free list blocks are always in address order:
$$addr(prev) < addr(curr) < addr(next)$$
 - » **con:** requires search
 - » **pro:** studies suggest fragmentation is lower than LIFO


Freeing With a LIFO Policy (Case 1)

Before



- Insert the freed block at the root of the list


After



Freeing With a LIFO Policy (Case 2)


conceptual graphic

Before



- Splice out predecessor block, coalesce both memory blocks, and insert the new block at the root of the list


After



Freeing With a LIFO Policy (Case 3)


conceptual graphic

Before



- Splice out successor block, coalesce both memory blocks and insert the new block at the root of the list


After



Freeing With a LIFO Policy (Case 4)


conceptual graphic

Before



- Splice out predecessor and successor blocks, coalesce all 3 memory blocks and insert the new block at the root of the list

After



Explicit List Summary

- **Comparison to implicit list:**
 - allocate is linear time in number of **free** blocks instead of **all** blocks
 - » **much faster** when most of the memory is full
 - slightly more complicated allocate and free since needs to splice blocks in and out of the list
 - some extra space for the links (2 extra words needed for each block)
- **Most common use of linked lists is in conjunction with segregated free lists**
 - keep multiple linked lists of different size classes, or possibly for different types of objects

Quiz 2


Assume that best-fit results in less external fragmentation than first-fit.

We are running an application with modest memory demands. Which allocation strategy is likely to result in better performance (in terms of time) for the application:


- a) best-fit**
- b) first-fit with LIFO insertion**
- c) first-fit with ordered insertion**

Keeping Track of Free Blocks

- Method 1: *implicit list* using length—links all blocks




- Method 2: *explicit list* among the free blocks using pointers



- Method 3: *segregated free list*
 - different free lists for different size classes
- Method 4: *blocks sorted by size*
 - can use a balanced tree (e.g. red-black tree) with pointers within each free block, and the length used as a key

Segregated List (Seglist) Allocators

- Each **size class** of blocks has its own free list



- Often have separate classes for each small size
- For larger sizes: One class for each two-power size


Seglist Allocator

- **Given an array of free lists, each one for some size class**
- **To allocate a block of size n :**
 - search appropriate free list for block of size $m > n$
 - if an appropriate block is found:
 - » split block and place fragment on appropriate list (optional)
 - if no block is found, try next larger class
 - repeat until block is found
- **If no block is found:**
 - request additional heap memory from OS (using `sbrk()`)
 - allocate block of n bytes from this new memory
 - place remainder as a single free block in largest size class

Seglist Allocator (cont.)

- **To free a block:**
 - coalesce and place on appropriate list
- **Advantages of seglist allocators**
 - higher throughput
 - » log time for power-of-two size classes
 - better memory utilization
 - » first-fit search of segregated free list approximates a best-fit search of entire heap.
 - » extreme case: giving each block its own size class is equivalent to best-fit

C vs. Storage Allocation




```
typedef struct block {  
    long size;  
    long payload[0];  
    long endsize;  
} block_t;
```



```
typedef struct free_block {  
    long size;  
    long payload[0];  
    struct free_block *next;  
    struct free_block *prev;  
    long endsize;  
} free_block_t;
```

Overloading Size



```
#define actual_size(s) ((s) & -2)  
#define allocated(s) ((s) & 1)
```

Is Previous Adjacent Block Free?

Size	?
Size	a
Next	
Prev	
Size	a

```
#define IsPrevAdjBlockFree(b) \
    !allocated(((long *) (b)) [-1])
```

```
static inline int  
IsPrevAdjBlockFree(free_block_t *b) {  
    long *prev_size = ((long *)b) [-1];  
    return !allocated(prev_size);  
}
```

```
static inline actual_size(long s) {  
    return s & -2;  
}
```

```
static inline allocated(long s) {  
    return s & 1;  
}
```

Is Next Adjacent Block Free?

Size	a
Next	
Prev	
Size	a
Size	?

```
#define IsNextAdjBlockFree(b) \
    !allocated(\
        ((long *) (b)) [actual_size(b->size)])

static inline int
IsNextAdjBlockFree(free_block_t *b) {
    long *next_size =
        ((long *) b) [actual_size(b->size)];
    return !allocated(next_size);
}
```